

Resources

Where to keep learning after this course ends — prompt libraries, prompt collections, cheat sheets, research papers, best practices hubs, and community examples.

DIFFICULTY LEVEL

All Levels

ESTIMATED READING TIME

35-45 minutes

ESTIMATED PRACTICE TIME

Ongoing

PREREQUISITES

Modules 01-14 (the entire bucket)

LEARNING OUTCOMES

By the end of this module, you will be able to: know where to find high-quality prompt examples and templates; read and extract value from AI/LLM research papers even without a research background; keep a personal, evolving best-practices reference; and stay current in a field that changes month to month.

Table of Contents

1.	Introduction	
2.	Before You Begin	
3.	Main Lesson	
3.1	Prompt Libraries	
3.2	Prompt Collections	
3.3	Cheat Sheets	
3.4	Research Papers	
3.5	Best Practices (Hubs)	
3.6	Community Examples	
4.	Visual Learning	
5.	Real World Applications	
6.	Practical Examples	
7.	Code Examples	
8.	Common Mistakes	
9.	Best Practices	
10.	Hands-on Activities	
11.	Mini Project	
12.	Review Questions	
13.	Cheat Sheet	
14.	Key Takeaways	
15.	Additional Resources & Course Wrap-Up	

1. Introduction

What is this topic?

This final module isn't a new technical skill — it's the map for continuing to learn after this course ends. AI engineering moves fast enough that no static course (including this one) can stay complete forever; this module is about knowing where to keep looking.

Why is it important?

The specific techniques in Modules 01-14 will keep evolving — new prompting patterns, new papers, new best practices. Knowing where to find reliable, current information is what turns "I finished a course once" into "I stay current in this field."

This Course (Snapshot)



Resource Habits



Continuous, Current Learning

Why should AI Engineers learn it?

Professionals in fast-moving fields are defined less by what they memorized once and more by their habits for staying current — the sources they check, the way they evaluate new claims, and how they keep a personal reference updated.

Where is it used in the real world?

Every practicing AI engineer maintains some version of this: bookmarked documentation, a personal prompt template library, a habit of skimming new papers, and a network of community sources they trust.

Why is it relevant today?

Given how fast this field moves, the resources module of any AI course is arguably the most important one for long-term relevance — techniques from Module 03 will get refined, new models will appear (Module 11), and this module is how you keep pace.

Note

This module points you toward resource *categories* and well-established, stable destinations rather than an exhaustive link list — specific URLs and sites change, but the categories of good resources stay consistent.

2. Before You Begin

RECAP

The whole bucket. This module assumes you've been through Modules 01-14 — it's designed to extend that foundation, not replace it.



Tip

Don't try to consume every resource category at once. Pick one (a research paper, a prompt library) and actually go look at it this week — habits form from small, real actions, not from reading a list.

3. Main Lesson

Each resource category follows: **Definition** → **Simple Explanation** → **Analogy** → **Where to Look** → **Why It Matters** → **Common Mistakes** → **Best Practices** → **Summary**.

3.1 Prompt Libraries

DEFINITION

A prompt library is a curated, often official collection of tested, reusable prompt templates for common tasks, usually organized by use case.

SIMPLE EXPLANATION

Rather than writing every prompt from scratch, prompt libraries give you a tested starting point for common tasks (summarization, classification, extraction) that you adapt to your specific need, applying Module 02's template principles.

 *Like a cookbook with tested recipes — you don't have to invent a lasagna recipe from first principles when a good, tested one already exists to adapt.*

WHERE TO LOOK

Anthropic's official prompt library (docs.claude.com) is a strong starting point, covering common use cases with tested example prompts for Claude specifically.

WHY IT MATTERS

Starting from a tested template is faster and more reliable than starting from a blank page, especially for common, well-understood task types.

COMMON MISTAKES

Using a library prompt verbatim without adapting it to your specific context, constraints, and output format needs.

BEST PRACTICES

Treat library prompts as a starting point, not a final answer — adapt and test against your own evaluation set (Module 09) before relying on them.

SUMMARY

Prompt libraries give you tested starting points for common tasks — a real time-saver when adapted thoughtfully rather than copied blindly.

3.2 Prompt Collections

DEFINITION

Prompt collections are broader, often community-maintained sets of prompts spanning many use cases and styles — typically larger and less curated/tested than an official prompt library.

SIMPLE EXPLANATION

Where a prompt library is usually official and quality-controlled, a prompt collection is often crowdsourced — useful for breadth and inspiration, but requiring more of your own judgment about quality.

 *Like the difference between a curated museum exhibit (library) and a huge public archive (collection) — the archive has more variety, but you do more of the sorting yourself.*

WHERE TO LOOK

GitHub "awesome-prompts" style repositories and community prompt-sharing sites are common sources — search for actively maintained ones with recent updates.

WHY IT MATTERS

Collections are useful for seeing a wide range of approaches to a problem, sparking ideas you might not have generated on your own.

COMMON MISTAKES

Trusting an unmaintained or unvetted collection's prompts as production-ready without testing them yourself.

BEST PRACTICES

Use collections for inspiration and breadth; always validate any prompt you actually adopt against your own test set (Module 09) before relying on it.

SUMMARY

Prompt collections trade curation for breadth — great for inspiration, but always verify before adopting into production use.

3.3 Cheat Sheets

DEFINITION

Cheat sheets are condensed, quick-reference summaries of key concepts, parameters, or syntax — exactly like the one-page cheat sheet at the end of every module in this course.

SIMPLE EXPLANATION

A good cheat sheet trades depth for speed of reference — you already understand the underlying concept (from a module like this one), and the cheat sheet just helps you recall it quickly without re-reading the full explanation.

 *Like a pilot's pre-flight checklist — not meant to teach you to fly, but to quickly remind you of key items once you already know what each one means.*

WHERE TO LOOK

Provider documentation (docs.claude.com) often includes quick-reference sections; this course's own Section 13 cheat sheets across all 15 modules form a personal reference set worth keeping.

WHY IT MATTERS

A good cheat sheet saves the time of re-deriving or re-reading something you've already learned once, but don't use often enough to have fully memorized.

COMMON MISTAKES

Relying on a cheat sheet to learn a concept for the first time, rather than as a memory aid for something already understood.

BEST PRACTICES

Build a cheat sheet only after you understand the underlying concept; keep it concise — a cheat sheet that's too long defeats its own purpose.

SUMMARY

Cheat sheets are fast-recall tools for already-understood concepts — this course's own Section 13s across all modules are a ready-made personal set.

3.4 Research Papers

DEFINITION

Research papers are the primary source of new AI/LLM techniques and findings — the original documents behind concepts covered throughout this course (e.g., Chain-of-Thought, Module 03; RAG, Module 07).

SIMPLE EXPLANATION

You don't need a research background to get value from papers — reading the abstract and conclusion first, then skimming figures/results, extracts most of the practical value without requiring you to follow every mathematical detail.

💡 *Like reading a movie's plot summary and reviews before deciding whether to watch the full 3-hour film — the abstract and conclusion give you 80% of the practical value in a fraction of the time.*

WHERE TO LOOK

arXiv.org (specifically the cs.CL and cs.AI categories) hosts most current LLM research, usually free and open; papers referenced throughout this course's "Additional Resources" sections are good starting points.

WHY IT MATTERS

Techniques you'll use in 6 months likely don't exist yet — reading papers (even just abstracts) is how you catch new developments before they become mainstream course content.

COMMON MISTAKES

Feeling you need to fully understand every mathematical detail before extracting practical value — abstract, intro, and conclusion often suffice for applied engineering purposes.

BEST PRACTICES

Read abstract → conclusion → skim figures/results first; only dive into full methodology if the paper is directly relevant to something you're building.

SUMMARY

Research papers are the original source of new techniques — approachable without a research background if you focus on abstract, conclusion, and results first.

🏠 3.5 Best Practices (Hubs)

DEFINITION

Best practices hubs are official, maintained guidance documents from model providers on how to use their specific models well — distinct from academic papers or community content.

SIMPLE EXPLANATION

Providers like Anthropic maintain living documentation specifically about prompting and using their models effectively — this guidance updates as models change, unlike a static course module or paper.

💡 *Like a car manufacturer's official maintenance guide versus a random online forum post — both might be useful, but the official guide is maintained specifically for your exact model and stays current with it.*

WHERE TO LOOK

docs.claude.com's prompt engineering section (referenced throughout this course) is the primary official best-practices hub for Claude specifically; other providers maintain equivalent documentation for their own models.

WHY IT MATTERS

Official guidance reflects the provider's own internal testing and model-specific quirks — often more precisely accurate for that specific model than general community wisdom.

COMMON MISTAKES

Applying generic "AI best practices" uniformly across every model family without checking whether the specific provider has model-specific guidance that differs.

BEST PRACTICES

Check official provider documentation first for model-specific guidance; treat it as more

authoritative than general community advice for that specific model.

SUMMARY

Best practices hubs are living, official, model-specific guidance — check them regularly since they update as models evolve.

👥3.6 Community Examples

DEFINITION

Community examples are real projects, code, and discussions shared by other practitioners — showing how concepts from this course get applied in practice, warts and all.

SIMPLE EXPLANATION

Official docs and papers show the idealized version of a technique; community examples show the messy reality — real edge cases, real debugging stories, real tradeoffs people made in production.

💡 *Like the difference between a cooking show's polished recipe demo and a home cook's forum post about what actually went wrong when they tried it — both valuable, in different ways.*

WHERE TO LOOK

GitHub repositories implementing real AI projects (many of Module 14's project types have open-source examples), developer forums, and provider community spaces.

WHY IT MATTERS

Community examples surface practical gotchas and real-world tradeoffs that polished official documentation sometimes glosses over.

COMMON MISTAKES

Treating any single community example as authoritative best practice without checking whether it's current, well-maintained, or widely validated by others.

BEST PRACTICES

Cross-reference community examples against official documentation; favor actively maintained, well-starred repositories over abandoned ones.

SUMMARY

Community examples show the messy, practical reality behind polished official guidance — valuable alongside, not instead of, official sources.

4. Visual Learning

📌 Resource Types by Authority vs. Currency

Resource Types

- **Official & Curated** *Prompt Libraries, Best Practices Hubs*
- **Rigorous but Slower-Moving** *Research Papers*
- **Broad & Fast-Moving** *Prompt Collections, Community Examples*

📖 How to Read a Research Paper Efficiently

- 1. Read the Abstract → What's the core claim?
- 2. Skip to the Conclusion → What did they actually find?
- 3. Skim Figures/Results → Does the data support the claim?
- 4. (Optional) Read Methods → Only if directly relevant to your work

📊 Resource Type Comparison

Resource Type	Best For	Update Frequency
Prompt Libraries	Tested starting templates	Periodic, provider-maintained
Prompt Collections	Breadth, inspiration	Variable, community-driven
Cheat Sheets	Fast recall of known concepts	Rarely (concepts are stable)
Research Papers	New techniques, rigor	Continuous (new papers daily)
Best Practices Hubs	Model-specific current guidance	Frequent, official
Community Examples	Real-world practical tradeoffs	Continuous, variable quality

5. Real World Applications

Scenario	Resource Type to Use
Starting a new classification task prompt	Prompt library (tested starting template)
Wondering if a new prompting technique exists for your problem	Research papers (arXiv)
Confirming Claude-specific prompting guidance	Anthropic's official best-practices docs
Seeing how others built a similar RAG system	Community examples (GitHub, forums)

6. Practical Examples

Beginner: Find one of each

Find one real example of each resource type from Section 3 (a prompt library entry, a research paper, a cheat sheet, etc.) and note where you found it.

Beginner: Practice the paper-reading method

Pick one paper referenced across this course's "Additional Resources" sections and apply Section 4's efficient-reading method (abstract → conclusion → figures).

Intermediate: Build your own cheat sheet

Pick your weakest module from this course and write your own one-page cheat sheet for it, without looking at the original — this both tests and reinforces your understanding.

Advanced: Evaluate a community example critically

Find a community project implementing one of Module 14's project types, and critically evaluate it: is it current? Does it follow the safety practices from Module 10? What would you change?

7. Code Examples

🔗Python — A personal prompt library starter (your own, growing resource)

```
# A simple starting point for YOUR OWN prompt library —
# the beginning of the habit this module is really about

PROMPT_LIBRARY = {
    "summarize": {
        "template": "Summarize the following in {n} sentences:\n\n{text}",
        "notes": "Works well with n=2-3. Add 'for a technical audience' for code docs.",
        "source": "Adapted from Module 02 templates",
    },
    "classify_sentiment": {
        "template": "Classify sentiment as Positive, Negative, or Neutral: {text}",
        "notes": "Add few-shot examples (Module 03) if consistency is inconsistent.",
        "source": "Module 03 few-shot example",
    },
    "extract_json": {
        "template": "Extract {fields} as JSON. Return ONLY valid JSON.\n\n{text}",
        "notes": "Always validate with Pydantic (Module 04) before using downstream.",
        "source": "Module 04 structured output pattern",
    },
}

def get_prompt(name, **kwargs):
    entry = PROMPT_LIBRARY[name]
    return entry["template"].format(**kwargs)

# Usage:
print(get_prompt("summarize", n=2, text="Long article text here..."))
```

📄JavaScript — A simple "resource check" habit tracker

```
// A tiny personal habit tracker for staying current —
// log one resource check per week to build the habit

const resourceLog = [];

function logResourceCheck(type, source, note) {
    resourceLog.push({ date: new Date().toISOString().split("T")[0], type, source, note });
    console.log(`Logged: [${type}] ${source} — ${note}`);
}

// Example usage over a few weeks:
logResourceCheck("paper", "arXiv: new RAG re-ranking approach", "Worth testing in Module 07 project");
logResourceCheck("docs", "docs.claude.com prompt engineering", "Checked for updates to XML tag guidance");
logResourceCheck("community", "GitHub: chat-with-pdf example", "Compared their chunking strategy to mine");
```



Tip

The real value of this section isn't the code itself — it's starting your own version of that `PROMPT_LIBRARY` dictionary today, and adding to it every time you write a prompt you'd want to reuse.

8. Common Mistakes

- **Treating this course as a permanent, complete reference** rather than a foundation to build on.
- **Using collection/community prompts verbatim** without adapting and testing them (Module 09) first.
- **Feeling you need full mathematical fluency** to get value from research papers — abstract/conclusion often suffice.
- **Applying generic best practices** without checking model-specific official guidance.
- **Trusting a single, possibly outdated community example** as authoritative without cross-referencing.
- **Never actually building a personal resource habit** — reading about resources without ever returning to them.

9. Best Practices

- **Keep a personal, growing prompt library** — even a simple file, as shown in Section 7.
- **Read papers efficiently:** abstract → conclusion → figures, diving deeper only when directly relevant.
- **Check official provider documentation** for model-specific guidance before applying generic advice.
- **Always test adopted prompts/techniques** against your own evaluation set (Module 09) before trusting them.
- **Build a small, regular habit** (a weekly resource check) rather than a one-time binge of reading.

10. Hands-on Activities

Easy 5 exercises

1. Bookmark one official resource from each of the 6 categories in Section 3.
2. Start your own prompt library file (even just 3 entries) following the Section 7 pattern.
3. Read one paper's abstract and conclusion (5 minutes) and summarize it in your own words.
4. Compare one of this course's cheat sheets against an official provider's quick-reference — note any differences.
5. Find one community example project related to something you built in Module 14.

Intermediate 3 exercises

1. Build out your personal prompt library to 10+ entries, each with notes on when/why to use it.
2. Read one full research paper referenced in this course and write a 3-sentence summary of its practical implications.
3. Set up the habit tracker from Section 7 and use it for 2 weeks, logging real resource checks.

Challenge 2 exercises

1. Find a claim from a community example that conflicts with official documentation, investigate which is more current/accurate, and write up what you found.
2. Identify one technique or development that has emerged since this course was written (search recent sources) and write a short module-style summary of it, following this course's own format.

11. Mini Project

Goal

Build your own personal "AI Engineering Reference" — a living document/repo combining a prompt library, a curated resource list, and cheat sheets from this entire course, that you'll actually keep using.

Requirements

- A place to store it (a GitHub repo, a personal notes app, or even a simple folder of files)

Step-by-Step Instructions

1. Create a "prompt_library" file/folder and add every reusable prompt template you've built across this course's 14 modules.
2. Create a "resources" file listing your personal top picks from each of Section 3's 6 categories, with a one-line note on why each is useful.
3. Copy or recreate the Section 13 cheat sheet from each of this course's 15 modules into one consolidated personal reference.
4. Set a recurring reminder (weekly or monthly) to add one new entry to any of these three files.
5. Six months from now, revisit this document and update anything that's changed.

Expected Output

```
my-ai-engineering-reference/  
├─ prompt_library.py (or .md)  — your reusable, tested prompts  
├─ resources.md                — your curated top picks per category  
├─ cheat_sheets/              — consolidated cheat sheets from all 15 modules  
└─ changelog.md               — dated notes on what you've updated and why
```

Possible Improvements

- Share it publicly (a GitHub repo) so others can benefit and you get community feedback.
- Add a simple search/tag system once your prompt library grows large.
- Cross-link entries back to the specific module/technique they came from, for future reference.

12. Review Questions

Multiple Choice (10)

1. A prompt library is best described as:

- A) A random unvetted collection B) A curated, often official set of tested prompt templates C) A type of model D) A research paper archive

2. Prompt collections differ from prompt libraries mainly by:

- A) Being more curated B) Being broader and typically less curated/vetted C) Having no prompts at all D) Being official only

3. A cheat sheet is best used as:

- A) A first-time learning tool for new concepts B) A fast-recall aid for concepts already understood C) A replacement for documentation D) A type of research paper

4. An efficient way to extract value from a research paper without deep expertise is:

- A) Reading only the title B) Abstract → conclusion → skim figures/results C) Memorizing every equation D) Skipping papers entirely

5. Best practices hubs from model providers are valuable because:

- A) They never update B) They reflect model-specific, current, official guidance C) They replace all other resources D) They're always wrong

6. Community examples are most valuable for:

- A) Official, guaranteed-correct guidance B) Real-world practical tradeoffs and gotchas C) Replacing research papers D) Nothing useful

7. Why should you test adopted prompts against your own evaluation set?

- A) It's unnecessary B) Library/collection prompts may not fit your specific context without adaptation C) It guarantees perfection D) It's required by every provider

8. arXiv is primarily a source for:

- A) Cheat sheets B) Research papers C) Community forums only D) Prompt libraries only

9. Why is a "resource habit" (regular, small checks) recommended over a one-time binge?

- A) It's slower overall B) The field moves continuously, so ongoing habits keep you current C) It's required by law D) One-time reading is always sufficient

10. Cross-referencing a community example against official documentation helps you:

- A) Waste time B) Verify currency and accuracy before trusting it C) Avoid all resources D) Guarantee it's outdated

Identification (5)

1. _____ is a curated, tested collection of reusable prompt templates, often official.

2. _____ is a condensed, quick-reference summary meant for fast recall of already-understood concepts.

3. _____ are the primary original source of new AI/LLM techniques and findings.

4. _____ are official, living guidance documents maintained by model providers.

5. _____ show real projects and practical tradeoffs shared by other practitioners.

True or False (5)

1. You need a full research background to get any practical value from a research paper.

2. Prompt libraries should be adapted and tested, not used verbatim without review.

3. Best practices hubs from providers update as models evolve.

4. This course, once finished, requires no further learning to stay current in AI engineering.

5. A small, regular resource-checking habit is more effective long-term than a

one-time information binge.

Situational (5)

- 1. You need a tested starting point for a summarization prompt. Which resource type would you check first?**
- 2. You want to know if there's a newer prompting technique than what's in this course. Where would you look?**
- 3. You found a community GitHub example that seems to contradict Anthropic's official documentation. What should you do?**
- 4. You keep forgetting Module 07's RAG pipeline steps months after finishing this course. What resource type helps most here?**
- 5. You want to build a habit of staying current without spending hours every week. What approach does this module recommend?**

Answer Key

Multiple Choice: 1-B, 2-B, 3-B, 4-B, 5-B, 6-B, 7-B, 8-B, 9-B, 10-B

Identification: 1-Prompt library, 2-Cheat sheet, 3-Research papers, 4-Best practices hubs, 5-Community examples

True or False: 1-False, 2-True, 3-True, 4-False, 5-True

Situational (sample reasoning):

1. A prompt library (e.g., Anthropic's official prompt library).
2. Research papers (arXiv) and provider best-practices hubs for the latest guidance.
3. Cross-reference both, favor official docs for authority, but check the docs' own update date too.
4. A cheat sheet — the fast-recall tool for a concept you already learned once.
5. A small, regular habit (e.g., a weekly resource check) rather than a one-time binge.

13. Cheat Sheet

 Libraries vs. Collections

Library — official, curated, tested.
Collection — broad, community, less vetted.

 Official vs. Community

Best practices hubs — official, model-specific, current.
Community examples — real-world, practical, variable quality.

 Reading Papers Fast

Abstract → Conclusion → Figures → (Methods if relevant).

 Golden Rule

Build small, regular resource habits — not one-time binges. Always test before trusting.

14. Key Takeaways

- Prompt libraries offer tested, official starting templates; prompt collections offer breadth at the cost of curation.
- Cheat sheets are fast-recall tools for concepts you already understand, not first-time learning material.
- Research papers are approachable without a research background — abstract, conclusion, and figures capture most practical value.
- Best practices hubs from providers are living, model-specific documents — check them for the most current, authoritative guidance.
- Community examples reveal real-world practical tradeoffs official docs sometimes omit — valuable alongside, not instead of, official sources.
- The highest-leverage habit from this module is small and regular: a weekly resource check beats a one-time binge, every time.

15. Additional Resources & Course Wrap-Up

Official Documentation

- Anthropic — docs.claude.com (prompt engineering guides, prompt library, best practices)
- arXiv.org — cs.CL and cs.AI categories for current LLM research

Communities

- Provider developer communities and forums for real-world discussion and troubleshooting.
- Actively-maintained GitHub repositories implementing the project types from Module 14.

🏆 You've completed the Prompt Engineering Mastery Bucket

Fifteen modules ago, this bucket started with the basics of what an LLM even is. It ends here, with you equipped to find, evaluate, and keep learning from the resources that will carry you past this course's own shelf life.

A quick recap of the journey: AI/LLM fundamentals (01) → prompting basics and techniques (02-03) → structured outputs and context engineering (04-05) → tool use and RAG (06-07) → agents (08) → optimization and safety (09-10) → the model/API landscape (11-12) → workflows (13) → real projects (14) → and now, resources to keep going (15).

The most valuable next step isn't reading another module — it's opening Module 14's capstone project (if you haven't already) and building something real with what you now know.