

Multimodal AI

Beyond text — vision models, image understanding and generation, OCR, speech-to-text and text-to-speech, audio and video understanding, multimodal prompting, document AI, and combining vision with RAG.

DIFFICULTY LEVEL

Intermediate

ESTIMATED READING TIME

55-70 minutes

ESTIMATED PRACTICE TIME

90-120 minutes

PREREQUISITES

Modules 01-19

LEARNING OUTCOMES

By the end of this module, you will be able to: send images to a vision-capable model and reason about their content; distinguish OCR from general image understanding; work with speech and audio inputs/outputs; understand video understanding's added complexity; write effective multimodal prompts; and combine vision with RAG for document-heavy use cases.

Table of Contents

1.	Introduction
2.	Before You Begin
3.	Main Lesson
3.1	Vision Models
3.2	Image Understanding
3.3	Image Generation
3.4	OCR
3.5	Speech-to-Text & Text-to-Speech
3.6	Audio Understanding
3.7	Video Understanding
3.8	Multimodal Prompting
3.9	Document AI
3.10	Vision + RAG
4.	Visual Learning
5.	Real World Applications
6.	Practical Examples
7.	Code Examples
8.	Common Mistakes
9.	Best Practices
10.	Hands-on Activities
11.	Mini Project
12.	Review Questions
13.	Cheat Sheet
14.	Key Takeaways
15.	Additional Resources

1. Introduction

What is this topic?

Every module so far has treated input and output as text. Multimodal AI extends everything you've learned — prompting (Module 02-03), structured output (Module 04), RAG (Module 07) — to images, audio, and video, alongside or instead of text.

Why is it important?

A huge share of real-world information isn't text: photos, scanned documents, screenshots, voice messages, video. Multimodal AI is what lets a system work directly with that information instead of requiring a human to manually transcribe or describe it first.



Why should AI Engineers learn it?

Vision-capable models are now standard across major providers (Module 11), and many real project types — Chat with PDF (Module 14, 3.5), Document AI, receipt/invoice processing — genuinely need to understand non-text content, not just describe it in a caption.

Where is it used in the real world?

Receipt and invoice scanning, medical image triage assistance, voice assistants, video content moderation, and accessibility tools (image descriptions for visually impaired users) all rely on multimodal AI.

Why is it relevant today?

Multimodal capability has moved from a specialized niche to a standard, expected feature across major model families (Gemini, Module 11, 3.3, in particular) — building genuinely useful AI products increasingly means working across modalities, not just text.

Note

Multimodal capability varies significantly by provider and model (Module 11) — always check current documentation for exactly which modalities a specific model supports natively versus via a separate specialized model.

2. Before You Begin

RECAP

Tokens (Module 01, 3.4) — images and audio also consume tokens (or an equivalent unit) once processed, affecting cost and context window the same way text does.

RECAP

RAG (Module 07) — Section 3.10 extends retrieval to work across visual content, not just text chunks.



Tip

Before reaching for a specialized single-purpose tool (a dedicated OCR library, a separate speech-to-text service), check whether your vision-capable LLM can already handle the task directly — it often can, with much less integration work.

3. Main Lesson

Each concept follows: **Definition** → **Simple Explanation** → **Analogy** → **Everyday Example** → **Why It Matters** → **Common Mistakes** → **Best Practices** → **Summary**.

👁️ 3.1 Vision Models

DEFINITION

Vision models are LLMs with the added capability to accept images as direct input alongside text, reasoning about visual content within the same conversation.

SIMPLE EXPLANATION

Rather than being a separate system, vision capability in modern LLMs (like Claude) is usually built into the same model you already use for text — you simply include an image alongside your text prompt in the same API call.

💡 *Like a colleague who can both read your written notes and look at a photo you hand them in the same conversation, reasoning about both together — not two separate people you'd have to coordinate between.*

EVERYDAY EXAMPLE

Sending a photo of a whiteboard sketch along with "turn this into a structured project plan," and getting back organized text that reasons about the actual visual layout of the sketch.

WHY IT MATTERS

Native vision capability means you can reason about images using all the same prompting techniques from Modules 02-03, without needing a completely separate computer-vision pipeline.

COMMON MISTAKES

Assuming every model in a family has vision capability — check provider documentation for which specific models support image input (Module 11).

BEST PRACTICES

Use native vision capability for most image-understanding tasks by default; reach for specialized computer-vision tools only for tasks requiring very specific, high-precision detection (e.g., exact pixel coordinates).

SUMMARY

Vision models bring image understanding into the same conversation and prompting toolkit you already use for text — a major simplification over separate CV pipelines.

🖼️ 3.2 Image Understanding

DEFINITION

Image understanding is a model's ability to describe, analyze, answer questions about, or reason over the actual content of an image.

SIMPLE EXPLANATION

This spans a spectrum from simple description ("what's in this image?") to detailed reasoning ("does this chart show an upward or downward trend, and by roughly how much?") — the same prompting principles from Module 02 (clear instructions, constraints) apply directly.

💡 *Like asking a sighted colleague to look at a diagram and answer specific questions about it, rather than just narrate everything they see — the quality of your question shapes the quality and focus of their answer.*

EVERYDAY EXAMPLE

Uploading a photo of a cluttered garage and asking "list the power tools visible in this image" — targeted image understanding, not just a general description.

WHY IT MATTERS

Precise, task-specific prompting (Module 02) applies directly to images — a vague "describe this" produces a vague answer, just as with text.

COMMON MISTAKES

Asking overly broad questions about complex images, then being surprised the model missed the specific detail you actually cared about.

BEST PRACTICES

Be as specific about what you want from an image as you would be with a text prompt — ask targeted questions rather than open-ended descriptions when you need something particular.

SUMMARY

Image understanding applies the same prompt-specificity principles from Module 02 to visual content — precision in, precision out.

3.3 Image Generation

DEFINITION

Image generation is producing new images from a text description, using a separate class of model (a diffusion model or similar) distinct from the text/vision LLMs covered elsewhere in this module.

SIMPLE EXPLANATION

Unlike image understanding (reading images), image generation is a fundamentally different technology (typically not the same LLM architecture) — most AI products combine an LLM for reasoning/planning with a separate specialized image-generation model for the actual creation step.

 *Like the difference between an art critic who can describe and analyze a painting (image understanding) and a painter who creates one from a verbal description (image generation) — related skills, but different specialists.*

EVERYDAY EXAMPLE

An LLM-powered design tool that uses the LLM to refine and clarify a user's vague request ("a friendly cartoon mascot") into a detailed prompt, then passes that prompt to a separate image-generation model to actually produce the image.

WHY IT MATTERS

Understanding that generation and understanding are typically different underlying technologies helps you architect systems correctly — don't expect a vision-understanding call to also generate images, or vice versa.

COMMON MISTAKES

Assuming the same model call that reads/understands images can also generate them, without checking whether the specific model/API actually supports generation.

BEST PRACTICES

Treat image generation as a distinct capability requiring its own model/API, often orchestrated by an LLM that refines the prompt first (Module 02's techniques apply to writing good generation prompts too).

SUMMARY

Image generation is a distinct capability from image understanding, typically requiring a separate specialized model — plan your architecture accordingly.

DEFINITION

OCR (Optical Character Recognition) is extracting text content from an image — turning pixels of text into actual, usable text characters.

SIMPLE EXPLANATION

OCR is a specific subset of image understanding (Section 3.2) focused purely on text extraction, rather than broader visual reasoning — modern vision-capable LLMs can often do this directly, reading text within an image as part of a normal prompt.

💡 *Like the specific skill of transcription versus general comprehension — OCR asks "what does this text say, exactly," a narrower and more precise task than "what's happening in this image."*

EVERYDAY EXAMPLE

Photographing a printed receipt and asking a vision-capable model to extract the total amount, date, and vendor name as structured JSON (Module 04) — OCR combined directly with structured output.

WHY IT MATTERS

OCR is foundational to Document AI (Section 3.9) and any workflow digitizing paper or image-based text — a hugely common real-world need.

COMMON MISTAKES

Not accounting for OCR accuracy dropping on low-quality images (blurry, low-resolution, poor lighting) — always validate extracted text for anything high-stakes.

BEST PRACTICES

Combine OCR with structured output validation (Module 04, 10) for anything downstream-consumed; provide the highest-quality image input reasonably available.

SUMMARY

OCR extracts text from images — a specific, high-value image understanding subtask that modern vision-capable LLMs handle directly, foundational for document processing.


🔊 3.5 Speech-to-Text & Text-to-Speech

DEFINITION

Speech-to-text (STT) converts spoken audio into written text; text-to-speech (TTS) converts written text into spoken audio — the two directions of the voice-text bridge.

SIMPLE EXPLANATION

These are typically handled by specialized models rather than the general-purpose LLM itself — a voice product usually chains STT → LLM (text reasoning) → TTS, similar in spirit to Module 03's prompt chaining pattern, but across modalities.



💡 *Like a translator working a phone call between two people speaking different languages — converting speech to text, having the actual conversation in text, then converting back to speech, seamlessly enough that it feels like one continuous exchange.*

EVERYDAY EXAMPLE

A voice assistant that transcribes a spoken question (STT), reasons about the answer with an LLM, then speaks the response back (TTS) — the classic voice-AI pipeline referenced in Module 18's latency discussions.

WHY IT MATTERS

Voice products are extremely latency-sensitive (Module 09, 3.6; Module 18, 3.6) — every

stage in the STT → LLM → TTS chain adds delay that compounds, making pipeline efficiency especially critical.

COMMON MISTAKES

Not accounting for cumulative latency across all 3 stages of the pipeline when designing a voice product's expected responsiveness.

BEST PRACTICES

Consider specialized low-latency inference platforms (Module 12, 3.5) for voice pipelines; use streaming (Module 18, 3.6) wherever the pipeline supports it to reduce perceived latency.

SUMMARY

STT and TTS bridge voice and text, typically chained around an LLM's text reasoning — a latency-sensitive pipeline worth designing carefully.

3.6 Audio Understanding

DEFINITION

Audio understanding is reasoning about audio content beyond just transcribing speech to text — tone, music, sound events, or non-speech audio characteristics.

SIMPLE EXPLANATION

This goes beyond Section 3.5's STT (which only extracts words) to actually reasoning about how something was said, what background sounds are present, or characteristics of non-speech audio like music.

 *Like the difference between a court stenographer who transcribes only the words spoken (STT) and a witness describing the tone of voice, background noise, and mood of a recording (audio understanding) — much richer information than the transcript alone.*

EVERYDAY EXAMPLE

Analyzing a customer service call recording not just for what was said (STT), but for detected frustration in the caller's tone — audio understanding beyond simple transcription.

WHY IT MATTERS

Some critical information (tone, urgency, emotional state, sound events) simply doesn't exist in a text transcript alone — audio understanding captures it directly.

COMMON MISTAKES

Relying only on a text transcript (STT) for tasks where tone or non-speech audio cues actually matter to the analysis.

BEST PRACTICES

Use audio understanding capability directly when tone, emotion, or non-speech audio content genuinely matters to the task, rather than working from transcript text alone.

SUMMARY

Audio understanding captures what a text transcript alone would miss — tone, sound events, and other non-verbal audio information.

3.7 Video Understanding

DEFINITION

Video understanding is reasoning about video content — combining visual frames over time with (often) audio, adding a temporal dimension beyond static image understanding (Section 3.2).

SIMPLE EXPLANATION

Video is the most complex modality covered in this module — it's essentially many images (frames) plus audio plus the added dimension of change and sequence over time, all of which the model needs to track together.

💡 *Like the difference between showing someone a single photograph of a soccer match (image understanding) versus the full match recording — the second requires tracking motion, sequence, and events unfolding over time, not just a frozen moment.*

EVERYDAY EXAMPLE

Asking a video-understanding model to summarize "what happened in this 2-minute security camera clip" — requiring it to track events, not just describe one static frame.

WHY IT MATTERS

Video processing is typically more token/compute-intensive than images (Module 01, 3.4) since it effectively involves many frames — cost and context window (Module 01, 3.6) considerations scale accordingly.

COMMON MISTAKES

Not accounting for video's higher token/cost footprint compared to a single image or text input when estimating costs (Module 09, 3.4).

BEST PRACTICES

Consider sampling key frames instead of full video where full temporal detail isn't needed; budget for video's higher token cost explicitly.

SUMMARY

Video understanding adds a temporal dimension on top of image understanding — powerful, but with meaningfully higher cost and complexity to plan for.

🗣️🖼️ 3.8 Multimodal Prompting

DEFINITION

Multimodal prompting is applying Module 02's prompt structure principles (instructions, context, constraints, Module 02, 3.1) to prompts that combine text with images, audio, or video.

SIMPLE EXPLANATION

The same structural discipline that makes text prompts effective — being specific, providing context, stating constraints — applies directly when an image or audio clip is part of the input; the multimodal content is simply additional context alongside your text instructions.

💡 *Like briefing someone on a task while also handing them a relevant photo — you still need to say clearly what you want them to do; the photo alone doesn't replace clear instructions.*

EVERYDAY EXAMPLE

"Here is a photo of a data table. Extract the values in the 'Revenue' column as a JSON array of numbers" combines Module 02's structure (clear instruction + output format, Module 04) with an image as context.

WHY IT MATTERS

A vague instruction accompanying an image ("what do you think?") produces vague results, just as it would with text alone — multimodal input doesn't reduce the need for prompt clarity.

COMMON MISTAKES

Assuming an image alone communicates intent clearly enough, and skimping on the accompanying text instructions.

BEST PRACTICES

Apply the same instruction/context/constraint/format structure from Module 02 to multimodal prompts, treating the non-text content as additional context, not a replacement for clear text instructions.

SUMMARY

Multimodal prompting is Module 02's structural discipline extended to non-text inputs — clarity still matters just as much.

📄 3.9 Document AI

DEFINITION

Document AI is the application of vision, OCR, and structured output together specifically for processing real-world documents — invoices, forms, contracts, receipts — often mixing text, tables, and layout.

SIMPLE EXPLANATION

This combines Section 3.4's OCR, Section 3.2's image understanding (for layout/structure), and Module 04's structured output into one focused application area, since real documents mix printed text, handwriting, tables, and visual layout that all carry meaning together.



Like a skilled office clerk who doesn't just read words off a form, but understands which field each piece of text belongs to based on its position and the form's layout — extraction that respects structure, not just raw text.

EVERYDAY EXAMPLE

Extracting line items, totals, and vendor details from a scanned invoice as validated structured JSON (Module 04), correctly associating each number with its corresponding table column despite varying layouts.

WHY IT MATTERS

Document AI is one of the highest-value, most common real-world multimodal applications — automating what would otherwise be tedious manual data entry across countless business processes.

COMMON MISTAKES

Treating document processing as pure OCR (text extraction only), missing the layout/structure understanding needed to correctly associate extracted values with their meaning.

BEST PRACTICES

Combine OCR, layout-aware image understanding, and strict schema validation (Module 04, 3.6) for reliable document processing; always validate extracted data given the high stakes of many document use cases (invoices, contracts).

SUMMARY

Document AI combines OCR, layout understanding, and structured output into one of the most practically valuable multimodal application areas.

📄 3.10 Vision + RAG

DEFINITION

Vision + RAG extends Module 07's retrieval-augmented generation to work with visual content — retrieving relevant images or image-derived text as context for a multimodal query.

SIMPLE EXPLANATION

This can mean two things: retrieving images themselves (via image embeddings) as relevant context, or extracting text from images first (Section 3.4's OCR) and feeding that

into a standard text-based RAG pipeline (Module 07) — the right approach depends on whether the visual layout itself matters.

💡 *Like a research assistant who, when helping with a visual archive, can either pull up the actual relevant photographs (image-based retrieval) or search through their written descriptions (text-based retrieval after OCR) — different approaches for different needs.*

EVERYDAY EXAMPLE

A product catalog assistant that retrieves the most visually similar product images for "find me something like this" queries, versus a scanned-document assistant that OCRs documents first and retrieves based on extracted text.

WHY IT MATTERS

Many real knowledge bases are image-heavy (scanned documents, product photos, diagrams) — pure text RAG (Module 07) misses this content entirely without a multimodal extension.

COMMON MISTAKES

Defaulting to OCR-then-text-RAG for every case, even when visual layout or appearance itself (not just the text) is what actually matters for retrieval relevance.

BEST PRACTICES

Choose image-embedding-based retrieval when visual similarity matters; choose OCR-then-text-RAG when the information is primarily textual but happens to be stored as an image (like a scanned document).

SUMMARY

Vision + RAG extends Module 07 to visual knowledge bases — the right retrieval approach depends on whether visual appearance or extracted text is what actually matters.

4. Visual Learning

📌 Modality Map

Multimodal AI

— Visual

Image Understanding, OCR, Video, Document AI

— Generative

Image Generation (separate model class)

— Audio

STT, TTS, Audio Understanding

🔄 The Voice Pipeline



📊 Modality Complexity & Cost

Modality	Relative Complexity	Relative Token/Cost Footprint
Text	Baseline	Baseline
Image	Moderate	Higher than text
Audio	Moderate-High	Depends on duration
Video	Highest	Highest (many frames + audio)

5. Real World Applications

Product	How these fundamentals show up
Expense/invoice management apps	Document AI extracts structured line items from photographed receipts.
Voice assistants	STT → LLM → TTS pipeline, latency-optimized (Module 18).
Accessibility tools	Image understanding generates descriptions for visually impaired users.
Visual product search	Vision + RAG retrieves visually similar catalog items.

6. Practical Examples

Beginner: Writing a multimodal prompt

Write a well-structured prompt (following Module 02's principles) for extracting specific information from a hypothetical photographed receipt.

Beginner: Choosing image vs. text RAG

For each: (1) a scanned PDF archive, (2) a product photo catalog — decide whether OCR-then-text-RAG or image-embedding RAG fits better, and why.

Intermediate: Estimating multimodal cost

Using Module 09's cost estimation approach, roughly compare the cost implications of processing 1,000 text queries vs. 1,000 image queries vs. 100 short videos.

Advanced: Designing a Document AI pipeline

Sketch a full pipeline (on paper) for processing scanned invoices: OCR → structure extraction → validation → storage, specifying what happens at each stage.

7. Code Examples

🐍Python — Image understanding + structured extraction (Document AI pattern)

```
# pip install anthropic
import anthropic
import base64
import json

client = anthropic.Anthropic(api_key="YOUR_API_KEY")

def extract_receipt_data(image_path):
    with open(image_path, "rb") as f:
        image_data = base64.standard_b64encode(f.read()).decode("utf-8")

    response = client.messages.create(
        model="claude-sonnet-4-6",
        max_tokens=300,
        messages=[{
            "role": "user",
            "content": [
                {
                    "type": "image",
                    "source": {"type": "base64", "media_type": "image/jpeg", "data": image_data}
                },
                {
                    "type": "text",
                    "text": ""Extract from this receipt as JSON: vendor, date, total.
Return ONLY valid JSON, no other text.""
                }
            ]
        }]
    )
    return json.loads(response.content[0].text)

data = extract_receipt_data("receipt.jpg")
print(data)  # {"vendor": "...", "date": "...", "total": "..."}

```

🌐JavaScript — Multimodal prompt with clear structure (Module 02 applied to images)

```
// npm install @anthropic-ai/sdk
import Anthropic from "@anthropic-ai/sdk";
const client = new Anthropic({ apiKey: process.env.ANTHROPIC_API_KEY });

async function analyzeChart(imageBase64) {
    const response = await client.messages.create({
        model: "claude-sonnet-4-6",
        max_tokens: 200,
        messages: [{
            role: "user",
            content: [
                { type: "image", source: { type: "base64", media_type: "image/png", data: imageBase64 } },
                {
                    type: "text",
                    // Module 02 structure: instruction + context + constraint + format, applied to an image
                    text: `You are a data analyst. Looking at this chart:
1. Identify the overall trend (up/down/flat).
2. State the approximate percentage change.
Constraint: answer in exactly 2 sentences.`
                }
            ]
        }],
    });
    return response.content[0].text;
}

```



Tip

Always validate extracted structured data from images (Section 7's JSON output) with the same Pydantic/Zod patterns from Module 04 — OCR and image understanding are not error-free, especially on lower-quality images.

8. Common Mistakes

- **Assuming every model in a family supports vision** without checking provider documentation.
- **Asking vague, open-ended questions about images** instead of specific, targeted ones.
- **Expecting one model call to both understand and generate images**, when these are typically separate capabilities.
- **Not validating OCR/extracted output**, especially from lower-quality images.
- **Ignoring cumulative latency** across a multi-stage voice pipeline (STT → LLM → TTS).
- **Underestimating video's token/cost footprint** compared to text or single images.
- **Defaulting to OCR-then-text-RAG** even when visual similarity itself is what actually matters.

9. Best Practices

- **Check current model documentation** for exact multimodal support before building around it.
- **Apply Module 02's prompt structure** to multimodal prompts — specificity matters just as much with images/audio.
- **Validate extracted structured data** from any visual/audio source, given real-world quality variance.
- **Budget explicitly for multimodal cost**, especially video's higher token footprint.
- **Choose image-embedding vs. OCR-then-text-RAG deliberately**, based on whether visual appearance or text content matters more.
- **Design voice pipelines for cumulative latency**, using streaming and low-latency platforms where needed.

10. Hands-on Activities

Easy 5 exercises

1. Send an image to a vision-capable model with a vague question, then again with a specific, targeted question — compare the answers.
2. List 3 real-world tasks that need OCR specifically, distinct from general image understanding.
3. Sketch the voice pipeline diagram from Section 4 from memory.
4. Explain, in your own words, why image generation typically needs a separate model from image understanding.
5. List 2 scenarios where video understanding's temporal dimension genuinely matters, versus 2 where a single frame would suffice.

Intermediate 3 exercises

1. Build the Python Document AI example from Section 7 with a real photographed document, and validate the output with Pydantic.
2. Write 3 different multimodal prompts (following Module 02's structure) for the same image, each requesting a different specific piece of information.
3. Design a cost estimate (Module 09-style) for a hypothetical product processing 500 receipt images per day.

Challenge 2 exercises

1. Design a full Vision + RAG architecture for a scanned-document archive, specifying OCR, chunking, embedding, and retrieval steps.
2. Design a voice assistant pipeline (STT → LLM → TTS) with specific latency targets for each stage, referencing Module 18's reliability techniques.

11. Mini Project

Goal

Build a small Document AI tool: photograph or find a sample document (receipt, invoice, or form), extract structured data via a vision-capable model, and validate the result.

Requirements

- An Anthropic API key
- Python with the Anthropic SDK and Pydantic
- A few sample document images (receipts, invoices, or forms)

Step-by-Step Instructions

1. Define a Pydantic model for the fields you want to extract (e.g., vendor, date, total, line items).
2. Build the image-input extraction function from Section 7, requesting exactly that JSON shape.
3. Run it against 3-5 sample document images of varying quality.
4. Validate each result against your Pydantic model, and note any that fail.
5. For any failures, investigate whether it was an image-quality issue or a prompt issue.

Expected Output

```
Document 1: ✓ Vendor='Cafe Luna', Date='2026-03-14', Total=12.50
Document 2: ✓ Vendor='Office Supply Co', Date='2026-03-15', Total=87.20
Document 3: ⚠ Validation failed — 'total' field missing
              (investigated: low image quality obscured the total line)
```

Possible Improvements

- Add line-item extraction (a nested list) for a more complete invoice-processing tool.
- Build a simple image-quality pre-check that flags likely-unreadable images before extraction.
- Combine with Module 07's RAG to make extracted documents searchable.

12. Review Questions

Multiple Choice (10)

1. Vision models are best described as:

- A) A separate product entirely B) LLMs with added capability to accept images as input C) Only for image generation D) Unrelated to text LLMs

2. Image generation typically requires:

- A) The same model as image understanding B) A separate, specialized model class C) No model at all D) Only text input

3. OCR is best described as:

- A) General image description B) Extracting text content specifically from an image C) Generating new images D) Audio transcription

4. A typical voice AI pipeline follows which order?

- A) TTS → LLM → STT B) STT → LLM → TTS C) LLM only, no STT/TTS D) Random order

5. Audio understanding goes beyond STT by also capturing:

- A) Nothing more B) Tone, sound events, and non-speech audio characteristics C) Only exact words D) Video frames

6. Video understanding adds which dimension beyond image understanding?

- A) Color B) Time/sequence C) Nothing new D) Text only

7. Multimodal prompting principles are:

- A) Completely different from text prompting B) The same structural principles from Module 02, applied to non-text input C) Unnecessary for images D) Only relevant for video

8. Document AI typically combines:

- A) Only OCR B) OCR, layout/image understanding, and structured output together C) Only image generation D) Only audio processing

9. When should you choose image-embedding-based retrieval over OCR-then-text-RAG?

- A) Never B) When visual similarity/appearance itself matters for relevance C) Always, regardless of use case D) Only for audio

10. Video typically has which cost/token characteristic compared to text or single images?

- A) Always cheaper B) Generally higher token/cost footprint C) Exactly the same D) Free

Identification (5)

- _____ extracts text content specifically from an image.
- _____ converts spoken audio into written text.
- _____ converts written text into spoken audio.
- _____ combines OCR, layout understanding, and structured output for real-world documents.
- _____ extends RAG to work with visual content, either via image embeddings or OCR-then-text.

True or False (5)

- All models within a given provider's family automatically support vision input.
- Image generation and image understanding are typically the same underlying model technology.
- Video understanding generally has a higher token/cost footprint than a single image.
- Multimodal prompts still benefit from clear, specific text instructions.
- Extracted data from images should be validated, similar to structured text output.

Situational (5)

1. You want to extract line items from a photographed invoice as validated JSON. What combination of techniques fits best?

2. You're building a product image search where visual similarity matters most. What retrieval approach fits?

3. Your voice assistant feels sluggish. What pipeline stages should you examine for latency?

4. You need to analyze customer call recordings for emotional tone, not just what was said. What capability do you need beyond STT?

5. Your image-extraction tool occasionally returns incomplete data on blurry photos. What should you add to your pipeline?

Answer Key

Multiple Choice: 1-B, 2-B, 3-B, 4-B, 5-B, 6-B, 7-B, 8-B, 9-B, 10-B

Identification: 1-OCR, 2-Speech-to-Text, 3-Text-to-Speech, 4-Document AI, 5-Vision + RAG

True or False: 1-False, 2-False, 3-True, 4-True, 5-True

Situational (sample reasoning):

1. Vision-based image understanding/OCR combined with structured output validation (Module 04).
2. Image-embedding-based retrieval rather than OCR-then-text-RAG.
3. Each stage of the STT → LLM → TTS pipeline, checking cumulative latency.
4. Audio understanding, which captures tone beyond just transcribed words.
5. Image-quality validation and/or a fallback/retry step, plus structured output validation to catch incomplete extractions.

13. Cheat Sheet

Understanding vs. Generation

Vision models **understand** images (same LLM).
Generation typically needs a **separate** model.

OCR vs. Document AI

OCR = text extraction. Document AI = OCR + layout + structured output.

Voice Pipeline

STT → LLM → TTS. Latency compounds across stages — optimize each one.

Vision + RAG

Visual similarity matters → image embeddings. Text matters → OCR then text-RAG.

14. Key Takeaways

- Vision models bring image understanding into the same LLM conversation and prompting toolkit as text.
- Image generation typically requires a separate, specialized model from image understanding.
- OCR is a focused text-extraction subtask; Document AI combines it with layout understanding and structured output.
- Voice pipelines (STT → LLM → TTS) are especially latency-sensitive, with delay compounding across stages.
- Audio and video understanding capture information (tone, sequence) that a transcript or single frame alone would miss.
- Multimodal prompting uses the same structural clarity principles as text prompting — non-text input doesn't reduce the need for specific instructions.
- Vision + RAG requires choosing between image-embedding retrieval and OCR-then-text-RAG based on what actually matters for relevance.

15. Additional Resources

Official Documentation

- Anthropic — vision capability documentation (docs.claude.com)
- Provider documentation for speech-to-text, text-to-speech, and image generation APIs (varies by provider, Module 12)

Related Modules

- Module 07 — RAG (extended to visual content in Section 3.10)
- Module 04 — Structured Outputs (essential for validating multimodal extraction results)
- Module 14 — Real Projects (Chat with PDF connects directly to Document AI)

Communities & Further Learning

- Provider changelogs for multimodal capability updates — this area evolves quickly across all major providers.